

Lab5: Let's Play GANs

Author: 311553007 應耀德

Lab5: Let's Play GANs

Introduction

Implementation details

Describe how you implement your model, including your choice of cGAN, model architectures, and loss functions.

DCGAN (based) + cGAN (condition)

DCGAN (based) + AC-GAN (condition)

Specify the hyperparameters (learning rate, epochs, etc.)

Results and Discussion

Show your results based on the testing data

Discuss the results of different models architectures

cGAN learning rate comparison

AC-GAN vs. cGAN

AC-GAN

cGAN

Introduction

這次 Lab 原先採用 DCGAN + AC-GAN 的組合進行訓練，發現不管是調整 learning rate, optimizer，採用 [ganhacks](#) 的建議，加入 Label smoothing, Dropouts to Discriminator 來使訓練 GAN 時穩定一點，其結果還是不盡理想，一直在 12% 左右震盪，最終 accuracy 皆落在 20% 左右。

後來改用 DCGAN + cGAN 的組合進行訓練，其結果真的好很多，最終 accuracy 落在 60% 左右。(隨便跑都這分數... AC-GAN 卻...)

Implementation details

Describe how you implement your model, including your choice of cGAN, model architectures, and loss functions.

DCGAN (based) + cGAN (condition)

Loss functions: Binary Cross-Entropy With Logits (BCEWithLogitsLoss) = BCELoss + Sigmoid

Objective function

$$\min_G \max_D V(D, G) = E_{x \sim p_{data}(x)} [\log D(x|c)] + E_{z \sim p_z(z)} [\log(1 - D(G(z|c)))]$$

Discriminator loss function

$$\min L_D = -E_{x \sim p_{data}(x)} [\log D(x|c)] + E_{x \sim p_g(x)} [\log D(x|c)]$$

Generator loss function

$$\min L_G = -E_{x \sim p_g(x)} [\log D(x|c)]$$

```
import torch
import torch.nn as nn
```

```
ngf = ndf = 64
```

```
class CGANGenerator(nn.Module):
    _in_channels: int

    def __init__(self, n_classes: int, z_dim: int, out_channels: int):
        super(CGANGenerator, self).__init__()

        self._z_dim = z_dim
        self._n_classes = n_classes

    def conv_block(
        in_channels: int,
        out_channels: int,
        kernel_size: int,
        stride: int,
        padding: int,
        bias: bool = False,
        normalization: bool = True,
    ):
        layers = [
            nn.ConvTranspose2d(
                in_channels,
                out_channels,
                kernel_size,
                stride,
                padding,
                bias=bias,
            )
        ]
        if normalization:
            layers.append(nn.BatchNorm2d(out_channels))
        layers.append(nn.LeakyReLU(inplace=True))
        return layers

    # like nn.Embedding but for multi-label task
    self.embedding = nn.Sequential(
        nn.ConvTranspose2d(
            n_classes, n_classes, 4, 1, 0, bias=False, groups=n_classes
        ),
        nn.LeakyReLU(inplace=True),
        nn.Flatten(),
        nn.Linear(n_classes * 4 * 4, z_dim, bias=False),
        nn.LeakyReLU(inplace=True),
    )

    self.main = nn.Sequential(
        *conv_block(z_dim * 2, ngf * 8, 4, 1, 0),
        # state size. (ngf*8) x 4 x 4
        *conv_block(ngf * 8, ngf * 4, 4, 2, 1),
        # state size. (ngf*4) x 8 x 8
        *conv_block(ngf * 4, ngf * 2, 4, 2, 1),
        # state size. (ngf*2) x 16 x 16
        *conv_block(ngf * 2, ngf, 4, 2, 1),
```

```

        # state size. (ngf) x 32 x 32
        nn.ConvTranspose2d(ngf, out_channels, 4, 2, 1, bias=False),
        nn.Tanh()
        # state size. (nc) x 64 x 64
    )

def forward(self, x: torch.Tensor, label: torch.Tensor):
    x = x.view(-1, self._z_dim, 1, 1)
    label = label.view(-1, self._n_classes, 1, 1)

    condition = self.embedding(label).view(-1, self._z_dim, 1, 1)
    return self.main(torch.concat((x, condition), dim=1))

class CGANDiscriminator(nn.Module):
    def __init__(self, in_channels: int, n_classes: int):
        super(CGANDiscriminator, self).__init__()

        self._image_size = 64
        self._in_channels = in_channels
        self._n_classes = n_classes

    def conv_block(
        in_channels: int,
        out_channels: int,
        kernel_size: int,
        stride: int,
        padding: int,
        bias: bool = False,
        normalization: bool = True,
    ):
        layers = [
            nn.Conv2d(
                in_channels,
                out_channels,
                kernel_size,
                stride,
                padding,
                bias=bias,
            )
        ]
        if normalization:
            layers.append(nn.BatchNorm2d(out_channels))
        layers.append(nn.LeakyReLU(0.2, inplace=True))
        return layers

    self.embedding = nn.Sequential(
        nn.ConvTranspose2d(
            n_classes, n_classes, 4, 1, 0, bias=False, groups=n_classes
        ),
        nn.LeakyReLU(0.2, inplace=True),
        nn.Flatten(),
        nn.Linear(
            n_classes * 4 * 4,
            in_channels * self._image_size**2,
            bias=False,

```

```

    ),
    nn.LeakyReLU(0.2, inplace=True),
)

self.convs = nn.Sequential(
    # input is (nc) x 64 x 64
    *conv_block(in_channels * 2, ndf, 4, 2, 1, normalization=False),
    # state size. (ndf) x 32 x 32
    *conv_block(ndf, ndf * 2, 4, 2, 1),
    # state size. (ndf*2) x 16 x 16
    *conv_block(ndf * 2, ndf * 4, 4, 2, 1),
    # state size. (ndf*4) x 8 x 8
    *conv_block(ndf * 4, ndf * 8, 4, 2, 1),
)

# output networks
# state size. (ndf*8) x 4 x 4
self.adversarial_network = nn.Sequential(
    nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False), nn.Flatten()
)

def forward(self, x: torch.Tensor, label: torch.Tensor):
    label = label.view(-1, self._n_classes, 1, 1)

    condition = self.embedding(label).view(
        -1, self._in_channels, self._image_size, self._image_size
    )
    output = self.convs(torch.concat([x, condition], dim=1))
    real_or_fake = self.adversarial_network(output)
    return real_or_fake

```

DCGAN (based) + AC-GAN (condition)

Loss functions: Binary Cross-Entropy With Logits (BCEWithLogitsLoss) = BCELoss + Sigmoid

Objective function

$$\begin{aligned}
 \max_D V(D) &= L_S + L_C \\
 \max_G V(G) &= -L_S + L_C \\
 L_S &= E_{x \sim p_{data}(x)} [\log D(x|c)] + E_{z \sim p_z(z)} [\log(1 - D(G(z|c)))] \\
 L_C &= E_{x \sim p_{data}(x)} [\log D(c|x)] + E_{z \sim p_z(z)} [\log D(c|G(z|c))]
 \end{aligned}$$

Discriminator loss function

$$\min L_D = -E_{x \sim p_{data}(x)} [\log D(x|c)] + E_{x \sim p_g(x)} [\log D(x|c)] + L_C$$

Generator loss function

$$\begin{aligned}
 \min L_G &= -E_{x \sim p_g(x)} [\log D(x|c)] + E_{z \sim p_z(z)} [\log D(c|G(z|c))] \\
 &= -E_{x \sim p_g(x)} [\log D(x|c)] + E_{x \sim p_g(x)} [\log D(c|x)]
 \end{aligned}$$

```

import torch
import torch.nn as nn

```

```
ngf = ndf = 64
```

```
class ACGANGenerator(nn.Module):
    _in_channels: int

    def __init__(self, n_classes: int, z_dim: int, out_channels: int):
        super(ACGANGenerator, self).__init__()

        self._z_dim = z_dim
        self._n_classes = n_classes

    def conv_block(
        in_channels: int,
        out_channels: int,
        kernel_size: int,
        stride: int,
        padding: int,
        bias: bool = False,
        normalization: bool = True,
    ):
        layers = [
            nn.ConvTranspose2d(
                in_channels,
                out_channels,
                kernel_size,
                stride,
                padding,
                bias=bias,
            ),
        ]
        if normalization:
            layers.append(nn.BatchNorm2d(out_channels))
        layers.append(nn.ReLU(inplace=True))
        return layers

    # state size. n_classes x 1
    # like nn.Embedding but for multi-label task
    self.embedding = nn.ConvTranspose2d(
        n_classes, n_classes, 4, 1, 0, bias=False, groups=n_classes
    )
    # state size. n_classes x 4 x 4

    self.l1 = nn.Sequential(
        *conv_block(z_dim, ngf * 8, 4, 1, 0, normalization=False)
    )
    self.main = nn.Sequential(
        # state size. (ngf*8) x 4 x 4
        *conv_block(ngf * 8 + n_classes, ngf * 4, 4, 2, 1),
        # state size. (ngf*4) x 8 x 8
        *conv_block(ngf * 4, ngf * 2, 4, 2, 1),
        # state size. (ngf*2) x 16 x 16
        *conv_block(ngf * 2, ngf, 4, 2, 1),
        # state size. (ngf) x 32 x 32
        nn.ConvTranspose2d(ngf, out_channels, 4, 2, 1, bias=False),
        nn.Tanh()
    )
```

```

        # state size. (nc) x 64 x 64
    )

def forward(self, x: torch.Tensor, label: torch.Tensor):
    x = x.view(-1, self._z_dim, 1, 1)
    label = label.view(-1, self._n_classes, 1, 1)

    x = self.l1(x)
    condition = self.embedding(label)
    return self.main(torch.concat((x, condition), dim=1))

class ACGANDiscriminator(nn.Module):
    def __init__(self, in_channels: int, n_classes: int):
        super(ACGANDiscriminator, self).__init__()

    def conv_block(
        in_channels: int,
        out_channels: int,
        kernel_size: int,
        stride: int,
        padding: int,
        bias: bool = False,
        normalization: bool = True,
    ):
        layers = [
            nn.Conv2d(
                in_channels,
                out_channels,
                kernel_size,
                stride,
                padding,
                bias=bias,
            ),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Dropout2d(0.5),
        ]
        if normalization:
            layers.append(nn.BatchNorm2d(out_channels))
        # layers.append(nn.LeakyReLU(0.2, inplace=True))
        return layers

    self.convs = nn.Sequential(
        # input is (nc) x 64 x 64
        *conv_block(in_channels, ndf, 4, 2, 1, normalization=False),
        # state size. (ndf) x 32 x 32
        *conv_block(ndf, ndf * 2, 4, 2, 1),
        # state size. (ndf*2) x 16 x 16
        *conv_block(ndf * 2, ndf * 4, 4, 2, 1),
        # state size. (ndf*4) x 8 x 8
        *conv_block(ndf * 4, ndf * 8, 4, 2, 1),
    )

    # output networks
    # state size. (ndf*8) x 4 x 4
    self.adversarial_network = nn.Sequential(

```

```

        nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False), nn.Flatten()
    )

    # state size. (ndf*8) x 4 x 4
    self.labels_classifier = nn.Sequential(
        nn.Conv2d(ndf * 8, n_classes, 4, 1, 0, bias=False), nn.Flatten()
    )

    def forward(self, x: torch.Tensor):
        output = self.convs(x)
        real_or_fake = self.adversarial_network(output)
        labels = self.labels_classifier(output)
        return real_or_fake, labels

```

Specify the hyperparameters (learning rate, epochs, etc.)

- Image transformation: `Resize(64) -> CenterCrop(64) -> ToTensor() -> Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))`
- Batch Size: 64
- Epoch Size: 300
- Z Dimension: 100
- Momentum: 0.5
- Generator
 - Optimizer: Adam
 - Learning Rate: $4e^{-4}$
- Discriminator
 - Optimizer: Adam
 - Learning Rate: $1e^{-4}$

Results and Discussion


Show your results based on the testing data

只列出成功訓練的結果 DCGAN + cGAN

產生的出來的 objects 都有類似情況，有的物體被切一半，有的像是兩個被切半的物體接在一起斜一邊

目前猜測是 Image transformation 的 CenterCrop 造成某些在邊緣的物體被切一半，需要另外實驗才能知道情況。

[Test accuracy: 56.9444%] [New Test accuracy: 63.0952%]




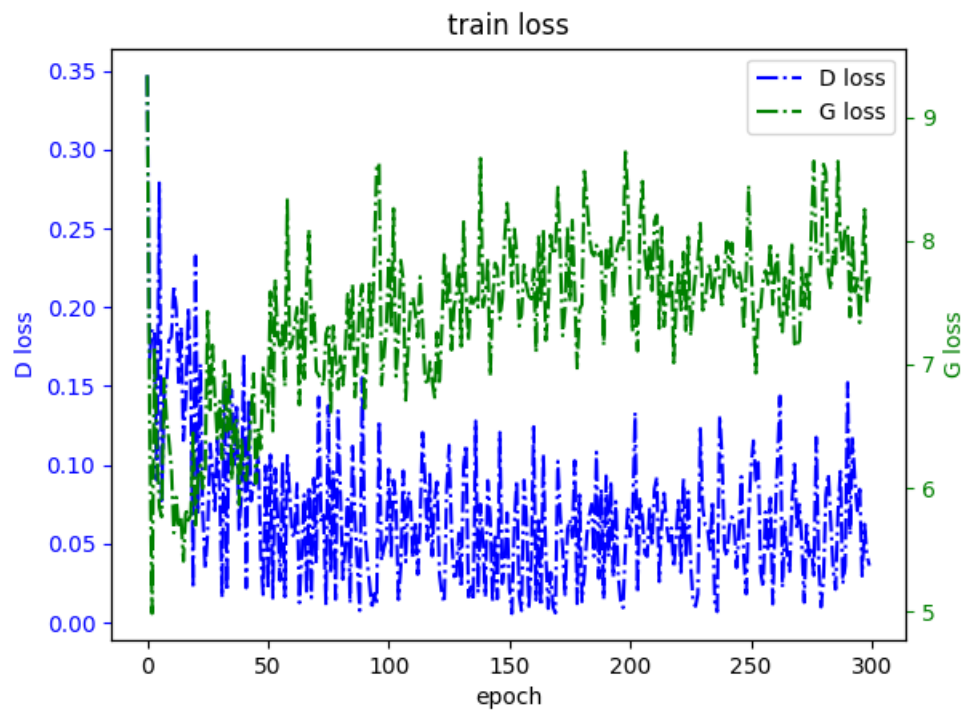
Discuss the results of different models architectures

cGAN learning rate comparison



discriminator 更新過快，「loss variance 過高，同時也牽制 discriminator 對 fake image (G loss) 的辨別能力，收斂速度較慢

反之，4:1 的 learning rate 較為穩定，也較早收斂

- Generator lr: $2e^{-4}$
Discriminator lr: $2e^{-4}$
Momentum: 0.5



- Generator lr: $4e^{-4}$
Discriminator lr: $1e^{-4}$
Momentum: 0.5




AC-GAN vs. cGAN


兩者架構除了 Discriminator 與 Loss function 不同以外，Generator 的 condition 方式皆採用 concat 的作法

下圖為 cGAN Discriminator 與 AC-GAN Discriminator 的對比圖，一個直接與 x 做 concat embedding，另一個則是把 condition 當作 classification target 來輔助 GAN 學習

**(b) cGANs,
hidden concat**
(Reed et al., 2016)



(c) AC-GANs
(Odena et al., 2017)





就如在 Introduction 中所說，自己認為是因為 AC-GAN 採用 auxiliary classifier 當作輔助訓練器，而在 GAN 前期 epoch 產生的圖片基本上都是雜訊，對 classifier 來說辨別效果有限甚至影響學習狀況，所以比 cGAN 這種直接 embed 進 image 後再丟入 Discriminator 中還難學習。

但因實驗時間太長，無法再繼續做相關實驗，只能等到課程結束後自己在慢慢 tuning。

以下是兩套架構所跑出的訓練情況，從 loss 訓練情況來看，很難判斷 learning 有問題，Generator loss 正常提高，代表它能產生 fake 假圖片資料，而不是一堆垃圾給 Discriminator，Discriminator 提高辨別精度 loss 降低，屬正常訓練情況。

但就從 accuracy 來看，差異就非常大大 (註: accuracy 單位為 epoch，是 average accuracy，與 test results 不一樣)

AC-GAN



cGAN

